

AI Financial Compliance & Transaction Security Platform

Complete English Project Explanation & Enterprise Assessment

Portfolio positioning: FinTech • Financial Security • AI Governance • Compliance • Auditability • Blockchain Integrity • Payment Infrastructure

Core principle

Let AI reason. Let deterministic software enforce. Let cryptography prove. Let humans authorize risk.

This document explains the project concept, architecture, security model, financial domain, AI boundary, compliance workflow, risk-based approval, audit design, blockchain role, payment abstraction, testing strategy, production gaps, and recommended roadmap.

1. Executive Summary

The project is a security-first financial control platform designed to process business invoices and financial transactions while keeping AI, compliance, authorization, payment execution, and auditability as separate responsibilities. The goal is not to build an AI chatbot that can move money. The goal is to build a controlled financial workflow in which AI can provide reasoning and recommendations, while deterministic software and explicit authorization policies enforce what the system is allowed to do.

The architecture combines Spring Boot, Spring Security, React/Redux, structured invoice ingestion, deterministic XML validation, local AI through Ollama, hybrid retrieval, compliance rules, risk scoring, human-in-the-loop approval, tokenization concepts, SHA-256 audit proofs, blockchain anchoring, and pluggable payment rails such as Stellar.

The project is strongest when positioned as an **AI Financial Compliance & Transaction Security Platform**, rather than as a blockchain demo or an invoice parser.

2. Problem Statement

Financial systems must balance automation with security, regulatory controls, auditability, privacy, and operational risk. AI can improve document understanding and compliance analysis, but an unrestricted AI component should not have direct authority over financial transactions.


This project addresses that problem by separating four concerns:


- **Reasoning:** local AI interprets structured and sanitized business context.
- **Enforcement:** deterministic validation and compliance rules decide what is permitted.
- **Authorization:** risk policies determine whether automation, human review, or blocking is required.
- **Proof:** cryptographic audit records and optional blockchain anchoring provide tamper-evident evidence.

3. Project Value Proposition

Dimension	Value
FinTech	Financial amounts, transactions, invoice processing, payment policies and payment-rail abstraction.
Security	Spring Security, security boundaries, tokenization concepts, log protection, fail-closed payment policy.
AI	Local LLM reasoning without giving the model direct authorization to move money.
Compliance	Deterministic validation combined with AI-assisted analysis and explicit risk decisions.
Audit	SHA-256 integrity proofs, append-only records and blockchain anchoring.
Enterprise	Separation of concerns, human oversight, configurable rules and future payment adapters.
Portfolio	Demonstrates Java/Spring, security, financial domain modeling, AI governance and distributed-system thinking.

4. High-Level Architecture





5. End-to-End Business Flow

- A business invoice enters through the secured API.
- The ingestion perimeter validates the document format and rejects unsafe XML constructs.
- ZUGFeRD/XRechnung-related structured data is parsed into controlled domain fields.
- Deterministic validation checks required fields and structural/business constraints before AI reasoning.
- Sensitive or unnecessary data is sanitized before being sent to the local AI boundary.
- The local LLM produces a structured compliance assessment rather than a direct payment instruction.
- Hybrid retrieval provides relevant compliance evidence and rules.
- The Compliance Engine combines deterministic validation and AI evidence.
- The Risk Engine assigns LOW, MEDIUM or HIGH risk.
- LOW-risk cases can be automatically approved; MEDIUM-risk cases require human approval; HIGH-risk cases are blocked.
- Approved decisions generate an audit record and cryptographic integrity proof.
- The hash can be anchored to a blockchain or other external integrity mechanism without publishing the invoice itself.
- Only an authorized payment workflow can call the configured PaymentRail.

6. Financial Domain Model

The financial core should remain independent from AI infrastructure. The current domain includes a MonetaryAmount value object and a FinancialTransaction abstraction. MonetaryAmount is based on BigDecimal, explicit currency information, scale and a defined rounding mode. FinancialTransaction contains transaction identity, amount and timestamp and exposes an ingestion validation perimeter.

```
MonetaryAmount
- BigDecimal amount
- Currency code
- scale
- rounding policy

FinancialTransaction
- transactionId
- amount
- timestamp
- validateIngestionPerimeter()

|
+---- FiatPaymentTransaction
+---- future StellarPaymentTransaction
+---- future SEPA / other transaction types
```

This separation is important because financial correctness should not depend on how an LLM behaves.

7. Invoice Ingestion & Deterministic Validation

The invoice edge is intentionally deterministic. XML should not be treated as trusted input merely because an AI model can read it.

Security requirements include protection against XML external entities and unsafe parser behavior. The parser should use secure DOM/parser settings and reject malformed or dangerous documents.

For a production implementation, the invoice layer should support profile-aware ZUGFeRD/XRechnung validation, including schema validation, business rules, code lists, VAT/tax consistency and document profile constraints. PDF/A-3 embedded XML extraction can be implemented as an additional adapter.

Important architectural rule: **schema and deterministic contract validation happen before AI reasoning.**

8. Local AI & AI Security

The local AI boundary is designed around Ollama or a similar locally hosted model. The motivation is to reduce unnecessary exposure of commercial invoice data and to keep model execution under the organization's infrastructure and security controls.

AI is advisory, not authoritative. A useful control is:

```
Invoice
-> sanitize
-> deterministic validation
-> structured context
-> local LLM
-> typed ComplianceAiAssessment
-> Compliance Engine
-> Risk Engine
-> Authorization
```


The AI layer should also enforce:

- prompt-injection detection and suspicious-instruction filtering;
- strict structured-output validation;
- model timeout and failure handling;
- model version and prompt version recording;
- confidence thresholds;

- no direct access to payment credentials;
- no unrestricted external network access;
- tool-level authorization if tools are introduced later.

9. Hybrid Retrieval


Compliance reasoning benefits from retrieval because rules, internal policies and evidence should be supplied as controlled context rather than relying only on model memory.



The current implementation is a lightweight retrieval implementation rather than a full production vector database. A production version can replace or augment it with a durable vector store while retaining deterministic filters, document versioning and effective dates.

10. Compliance Engine

The Compliance Engine is one of the central design elements. It should not reduce compliance to a single LLM answer. Instead, it should combine deterministic validation, retrieved evidence, configured rules and AI interpretation.



Rules should eventually be versioned and associated with effective dates and authoritative sources. This makes decisions reproducible and auditable.

11. Risk Engine & Human-in-the-Loop

Risk-based automation is a key enterprise feature. Not every transaction should receive the same treatment.

Risk	Default action	Rationale
LOW	Auto approve	Deterministic checks pass and risk is within configured limits.
MEDIUM	Human review	Requires explicit authorization or additional evidence.
HIGH	Block	Potential fraud, policy violation, hash mismatch, dangerous input or other critical condition.

Additional controls can force human review for high transaction amounts, unusual counterparties, tax inconsistencies, missing evidence or cryptographic mismatches.

12. Payment Security & Tokenization

Payment data should be isolated from the rest of the business domain. A tokenization boundary can convert sensitive card data into an application token before the core system processes the transaction.

```
Raw payment credential
|
v
Tokenization Edge
|
v
Token ID
|
+---- Core financial services
+---- Logs / monitoring
+---- compliance workflows
```

A critical regression test should verify that the raw PAN never appears in application logs. The current tokenization component is a demonstration and is explicitly not a PCI-certified vault or HSM implementation.

For production, secrets and payment cryptographic material should be protected by appropriate HSM/KMS/Vault infrastructure and the organization's PCI/security architecture.

13. Payment Authorization Security

The most important security invariant is that a compliance failure cannot silently turn into a payment.

```
Compliance = DENY
|
+----> PaymentRail MUST NOT be called

Compliance = REVIEW
|
+----> humanApproved == true
|
+----> PaymentRail

Compliance = APPROVE
|
+----> risk/policy checks
|
+----> PaymentRail
```

This is stronger than relying on an AI prompt such as 'do not pay risky invoices'. The underlying service layer must enforce the rule.

14. Auditability & Cryptographic Integrity

The audit system should capture what happened, who or what initiated it, which rule and model versions were involved, and what decision was made.

```
AuditEvent
eventId
actor
timestamp
action
resourceId
decision
riskLevel
ruleVersion
modelVersion
previousHash
currentHash
```

A hash chain provides tamper-evident ordering:

```
Event 1 -> Hash 1
Event 2 + Hash 1 -> Hash 2
Event 3 + Hash 2 -> Hash 3
```

An external blockchain can then anchor a selected audit hash. The invoice itself, personal data, PAN, IBAN or commercial secrets should not be written to a public blockchain.

Blockchain anchoring is an integrity mechanism. It should not be presented by itself as proof of regulatory compliance or legal/recordkeeping compliance.

15. Blockchain & Web3 Role

Web3j is most useful here as an adapter for an integrity layer, not as a reason to put all financial data on-chain. Stellar can remain the payment/settlement adapter while an Ethereum-compatible network can optionally provide an external audit anchor.

```
PaymentRail
|
+---- StellarPaymentAdapter
|
+---- future SEPA adapter
|
+---- other rails

AuditAnchor
|
+---- local cryptographic ledger
|
+---- Web3j / EVM anchor
+---- future enterprise ledger
```

This separation prevents the project from becoming 'blockchain for blockchain's sake'.

16. Security Architecture

```
External Input
|
v
[Spring Security]
|
v
[Input / XML Security]
|
v
[Sanitization + DLP boundary]
|
v
[Local AI Guardrails]
|
v
[Deterministic Compliance]
|
v
[Risk / Authorization]
|
+----- BLOCK
|
+----- HUMAN REVIEW
|
+----- APPROVE
|
v
[Audit Integrity]
|
v
[PaymentRail]
```

Security is therefore layered. The AI component is not the security boundary by itself; the financial service layer remains authoritative.

17. API Design

The project can expose separate APIs for invoices, approvals, payments and audit data. Example responsibilities:

Endpoint concept	Purpose
POST /api/invoices	Upload and process invoice data.
POST /api/approvals	Record explicit human approval.
POST /api/payments	Execute an authorized payment request.
GET /api/audit/...	Retrieve audit evidence and integrity status.

Production APIs should add OAuth2/OIDC, RBAC/ABAC, rate limiting, request IDs, idempotency keys, validation, structured error responses and audit correlation IDs.

18. Observability

An enterprise deployment needs more than application logs. Recommended observability includes:

- structured JSON logs without sensitive payment data;
- correlation/request IDs;
- metrics for invoice processing, compliance decisions, AI latency and payment failures;
- traces across ingestion → AI → compliance → authorization → payment;
- security alerts for repeated failures or suspicious input;
- health/readiness endpoints;
- model and rule version metrics.

19. Testing Strategy

Test	What it proves
FinancialPrecisionTest	Money uses BigDecimal and defined scale/rounding semantics.
ArchitectureTest	Financial/security boundaries and dependency rules are respected.
XmlSecurityTest	Unsafe XML constructs are rejected.
TokenizationRegressionTest	Sensitive payment data is converted to a token.
LogComplianceTest	Raw PAN is not present in logs.
PromptInjectionTest	Suspicious instructions do not cross the AI boundary.
ComplianceRegressionTest	Invalid, missing and valid cases produce expected decisions.
RiskEngineTest	Risk classification is deterministic for configured policies.
PaymentAuthorizationTest	DENY never calls PaymentRail; REVIEW requires human approval.
AuditIntegrityTest	Hash chain detects tampering.
Integration tests	Spring Security, persistence and service boundaries work together.

The strongest tests are invariant tests: they prove that dangerous states cannot reach the payment execution boundary.

20. Enterprise Assessment

Area	Assessm ent	Comment
FinTech domain design	8.5/10	Strong financial model, invoice and payment concepts.
Application security	8/10	Good layered security direction.
Payment security	7.5/10	Good architecture; vault/HSM remains production work.
AI security	7.5/10	Local AI and controlled boundary are strong; more controls needed.
Compliance architecture	8.5/10	Deterministic + AI + risk + human review is compelling.
Blockchain/Web3	8/10	Strong as integrity/audit layer, not as a generic blockchain demo.
Auditability	8.5/10	Hash chain and evidence model add significant value.
Architecture	9/10	Clear separation of security, compliance, risk, approval and payment.
Production readiness	5.5/10	Persistence, IAM, HSM/KMS, observability and deployment hardening remain.
Testing	6.5/10	Good foundation; needs broader integration and security coverage.
Enterprise relevance	9/10	Highly relevant to payments, ERP, banking and compliance workflows.
Portfolio value	9/10	Demonstrates senior-level cross-domain engineering potential.

21. Current Limitations

- The current retrieval implementation is lightweight and should not be described as a production vector database.
- Invoice support is an MVP parser/validator rather than a complete implementation of every ZUGFeRD/XRechnung profile and PDF/A-3 workflow.
- The tokenization component is a demonstration and is not a certified PCI vault.
- The Stellar adapter is simulation-first and does not represent production signing and transaction submission.
- Audit and approval storage still need durable persistence for a production deployment.
- The AI output contract should be strengthened with strict schema validation and model/prompt versioning.
- Full compilation, dependency resolution and containerized integration testing should be performed in a Maven/Docker-capable CI environment.

22. Production Roadmap

Phase 1 – Secure MVP
Invoice ingestion
XML security
BigDecimal domain
Local Ollama
Compliance + risk
Payment policy

Audit hash chain

Phase 2 – Enterprise Controls

PostgreSQL
OAuth2/OIDC
RBAC/ABAC
Vault/KMS
durable approvals
structured audit events
OpenTelemetry

Phase 3 – Financial Integrations

complete ZUGFeRD/XRechnung profiles
SEPA adapter
production Stellar integration
idempotency
reconciliation

Phase 4 – Integrity Infrastructure

Web3j/EVM audit anchoring
independent verification service
key rotation
signed audit manifests

Phase 5 – AI Governance

model registry
prompt versioning
evaluation datasets
adversarial testing
model risk controls
explainability and human override

23. Recommended Repository Structure

backend/
security/
invoice/
ai/
retrieval/
compliance/
risk/
approval/
tokenization/
transaction/
payment/
audit/
ledger/
config/
controller/
model/

frontend/
components/
features/
store/
services/

docs/
ARCHITECTURE.md
SECURITY_ASSESSMENT.md
TEST_STRATEGY.md
THREAT_MODEL.md
API.md

examples/
invoice.xml

24. Threat Model

Threat	Primary control
Malicious XML	Secure parser, schema validation, XXE protection.
Prompt injection	Input sanitization, instruction filtering, constrained AI context.
AI hallucination	Deterministic rules, structured output and human review.
Unauthorized payment	Payment policy, authorization service and fail-closed behavior.
PAN leakage	Tokenization boundary and log regression tests.
Audit tampering	Hash chain and optional external blockchain anchor.
Secret compromise	Production KMS/HSM/Vault and key rotation.
Replay / duplicate payment	Idempotency keys and transaction state machine.
External AI data leakage	Local model, network isolation and controlled egress.

25. Key Security Invariants

- AI must never have direct authorization to move money.
- DENY must never call PaymentRail.
- REVIEW must require explicit human authorization.
- Raw PAN must never appear in logs.
- Invoice content must never be written to a public blockchain.
- Deterministic schema/security validation must occur before AI reasoning.
- Payment secrets must not be stored in source code or ordinary application configuration.
- Audit records must be integrity-protected and reproducible.

26. Definition of Done

- A malicious XML document is rejected safely.
- A valid invoice can be parsed into typed financial data.
- Financial calculations use explicit decimal semantics.
- Local AI can assess controlled context and return a validated structured result.
- Compliance decisions are reproducible from rules, evidence and recorded versions.
- Risk determines automatic approval, human review or blocking.
- A blocked transaction cannot reach the payment adapter.
- A reviewed transaction cannot reach payment without explicit approval.
- Sensitive payment credentials are tokenized before entering the core workflow.
- Sensitive credentials do not appear in logs.
- Every financial decision creates an audit record.
- Audit records have verifiable cryptographic integrity.
- Blockchain anchoring, when enabled, stores only appropriate integrity metadata.

- The system has automated unit, architecture, security and integration tests.
- Production deployment uses proper identity, secrets management, persistence and observability.

27. Portfolio Positioning

For a technical interview or portfolio, the project should be presented as a controlled financial platform rather than as a collection of technologies. The most important engineering story is the set of boundaries between AI, security, compliance and payment execution.

```
AI = reasoning
Rules = enforcement
Risk = decision policy
Human = authorization for exceptions
Cryptography = integrity proof
Blockchain = external anchor
PaymentRail = controlled execution
```

This makes the project relevant to roles involving Java/Spring, FinTech, payments, application security, AI infrastructure, compliance automation and enterprise architecture.

28. Final Assessment

Architectural concept: 9/10

The separation of reasoning, enforcement, authorization and proof is strong and professionally relevant.

Current implementation maturity: approximately 6–6.5/10

The architecture is ahead of the production implementation. Persistence, IAM, HSM/KMS, full invoice profile support, observability, stronger AI governance and production payment integrations are the main gaps.

FinTech relevance: 8.5/10

The project contains meaningful financial domain modeling, invoice processing, compliance, risk and payment concepts.

Security relevance: 8/10

The layered security approach and payment authorization invariants are strong portfolio material.

Enterprise potential: 9/10

With production hardening, the architecture could map naturally to banking, payment, ERP and financial compliance systems.

Portfolio value: 9/10

The project demonstrates cross-functional engineering instead of a single-framework demo.

Final architecture statement

Let AI reason. Let deterministic software enforce. Let cryptography prove. Let humans authorize risk.

That principle is the strongest single sentence to use when explaining the project to a senior engineer, security architect, FinTech recruiter or hiring manager.